(a) World model, imagined rollout, and the monitor

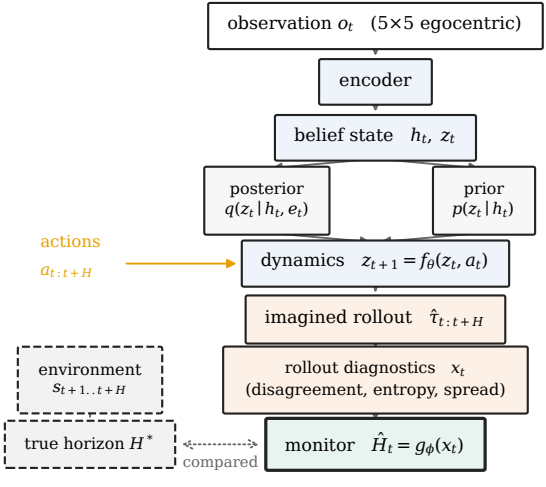

(b) Three world-model architectures

M1 GRU-RSSM (large)

deter 512 / stoch 48

recurrence

M2 GRU-RSSM (small)

deter 256 / stoch 16

recurrence

M3 Transformer

d 256, 2 layers, 4 heads

causal attention

identical data, seeds, targets and downstream code; only dynamics differ